


Guanghai Qin

Curriculum Vitae

Updated May 26, 2024

CONTACT

Department of Computer Science
Johns Hopkins University
3400 N Charles Street
Baltimore, MD 21218, USA

Email: gqin2@jhu.edu
Website: <https://gqin.me>
ORCID: 0000-0002-3009-8614 

EDUCATION

Johns Hopkins University
M.S.E. in Computer Science
Ph.D. in Computer Science
Mentor: [Prof. Benjamin Van Durme](#)

Maryland, USA
Aug 2019 – May 2023
Aug 2019 – Present

Peking University
B.S. in Physics (primary major)
B.S. in Computer Science (secondary major)
Physics GPA: 3.52; CS GPA: 3.72.

Beijing, P. R. China
Sept 2015 – Jun 2019
Sept 2015 – Jun 2019

PROFESSIONAL EXPERIENCE

Microsoft Research Lab (MSR)
Research Intern. Mentor: [Corby Rosset](#)

May 2023 – Aug 2023

I researched efficient methods for long-context LLMs. I proposed a method to compress the context of LLaMA with a compression ratio of up to 20x with minimal performance tradeoffs.

Microsoft Semantic Machines
Research Intern. Mentor: [Anthony Platanios](#)

May 2022 – Aug 2022

I studied a new research problem for user action predictions. I built a 2TiB dataset from GitHub and implemented a model based on graph neural networks (GNNs) to predict user actions (e.g. issues).

Johns Hopkins University
Visiting Scholar. Mentor: [Hongyuan Mei](#) and [Jason Einser](#).

Jun 2018 – Nov 2018

I worked on temporal event stream modeling. We proposed a particle smoothing solution to sample events from a neural Hawkes process, which may interact with a deductive temporal database such as Datalog.

Microsoft Research-Asia (MSRA)
Research Intern. Mentor: [Jin-Ge Yao](#) and [Chin-Yew Lin](#).

Nov 2017 – Jun 2018

I proposed a Semi-HMMs-based statistical model for grounding natural language to structured data, which can be used to induce templates for data-to-text generation.

HONORS

Best Short Paper Awardee in NAACL	2021
Outstanding Reviewer in EMNLP	2019
Award for Scientific Research (Peking University)	2018

May 4th Scholarship (Peking University)	2016 and 2018
Merit Student (Peking University)	2017
Award for Academic Diligence (Peking University)	2016
Chinese Physics Olympiad (CPhO) Silver Medalist	2014

PROFESSIONAL ACTIVITIES

Reviewer for NeurIPS (2019, 2020, 2021, 2022, and 2023)
 Reviewer for ICLR (2019, 2020, 2021, 2023, and 2024)
 Reviewer for ACL (2021)
 Reviewer for EMNLP (2020, 2021, and 2022)
 Reviewer for NAACL (2024)
 Reviewer for AACL (2021)

SKILLS

Programming languages: Python, Rust, JAVA, C/C++, Shell, \LaTeX , and SQL.
Machine learning tools: PyTorch, Lightning AI, DeepSpeed, and FAISS.
Network/Web/Database: Front- and back-end. MySQL database.

TEACHING EXPERIENCE

TA for *Machine Learning* (CS 601.475) at Johns Hopkins University in 2022.

PREPRINTS

2. Corby Rosset, Ho-Lam Chung, **Guanghui Qin**, Ethan C Chau, Zhuo Feng, Ahmed Awadallah, Jennifer Neville, and Nikhil Rao. [Researchy Questions: A Dataset of Multi-Perspective, Decompositional Questions for LLM Web Agents](#). In *arXiv*. 2024.
1. Wenting Tan, Yunmo Chen, Tongnfei Chen, **Guanghui Qin**, Haoran Xu, Heidi C Zhang, Benjamin Van Durme, and Phillip Koehn. [Streaming Sequence Transduction through Dynamic Compression](#). In *arXiv*. 2024.

PUBLICATIONS

13. **Guanghui Qin**, Corby Rosset, Ethan C Chau, Nikhil Rao, and Benjamin Van Durme. [Dodo: Dynamic Contextual Compression for Decoder-only LMs](#). In *Proceedings of Annual Meeting of the Association for Computational Linguistics (ACL, main)*. 2024.
12. Yiyang Lin*, Dhiman S Pal*, Parijat Banerjee, Tatsat Banerjee, **Guanghui Qin**, Yu Deng, Jane Borleis, Pablo A Iglesias, and Peter Devreotes. [Ras suppression potentiates actomyosin-driven polarization in migrating cells](#). In *Nature Cell Biology (accepted, issue number TBA)*. 2024.
11. **Guanghui Qin** and Benjamin Van Durme. [Nugget: Neural Agglomerative Embeddings of Text](#). In *Proceedings of the Conference on International Conference on Machine Learning (ICML)*. 2023.
10. **Guanghui Qin**, Yukun Feng, and Benjamin Van Durme. [The NLP Task Effectiveness of Long-Range Transformers](#). In *Proceedings of the European Chapter of the Association for Computational Linguistics (EACL, oral)*. 2023.
9. **Guanghui Qin** and Jason Eisner. [Learning How to Ask: Querying LMs with Mixtures of Soft Prompts](#). In *Proceedings of Conference of the North American Chapter of the Association for Computational Linguistics (NAACL, short, oral)*. 2021.

Best Short Paper Award

8. Mahsa Yarmohammadi, Shijie Wu, Marc Marone, Haoran Xu, Seth Ebner, **Guanghui Qin**, Yunmo Chen, Jialiang Guo, Craig Harman, Kenton Murray, Aaron S White, Mark Dredze, and Benjamin Van Durme. [Everything Is All It Takes: A Multipronged Strategy for Zero-Shot Cross-Lingual Information Extraction](#). In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP, oral)*. 2021.
7. Patrick Xia*, **Guanghui Qin***, Siddharth Vashishtha, Yunmo Chen, Tongfei Chen, Chandler May, Craig Harman, Kyle Rawlins, Aaron Steven White, and Benjamin Van Durme. [LOME: Large Ontology Multilingual Extraction](#). In *Proceedings of Conference of the European Chapter of the Association for Computational Linguistics (EACL, demo)*. 2021.
6. Ryan Culkin, J Edward Hu, Elias Stengel-Eskin, **Guanghui Qin**, and Benjamin Van Durme. [Iterative Paraphrastic Augmentation with Discriminative Span-based Alignment](#). In *Transactions of the Association for Computational Linguistics (TACL)*, 9:494-509. 2021.
5. Hongyuan Mei, **Guanghui Qin**, Minjie Xu, and Jason Eisner. [Neural Datalog Through Time: Informed Temporal Modeling via Logical Specification](#). In *Proceedings of the Conference on International Conference on Machine Learning (ICML, oral)*. 2020.
4. Singh Abhinav, Patrick Xia, **Guanghui Qin**, Mahsa Yarmohammadi, and Benjamin Van Durme. [CopyNext: Explicit Span Copying and Alignment in Sequence to Sequence Models](#). In *Proceedings of Workshop on Structured Prediction for NLP*. 2020.
3. Hongyuan Mei, **Guanghui Qin**, and Jason Eisner. [Imputing Missing Events in Continuous-Time Event Streams](#). In *Proceedings of the Conference on International Conference on Machine Learning (ICML, oral)*. 2019.
2. **Guanghui Qin**, Jin-Ge Yao, Xuening Wang, Jinpeng Wang, and Chin-Yew Lin. [Learning Latent Semantic Annotations for Grounding Natural Language to Structured Data](#). In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP, oral)*. 2018.
1. Longxu Dou, **Guanghui Qin**, Jinpeng Wang, Jin-Ge Yao, and Chin-Yew Lin. [Data2Text Studio: Automated Text Generation from Structured Data](#). In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP, demo)*. 2018.

Asterisk * indicates equal contribution

CITATIONS

Visit my Google Scholar profile to view my citations:

<https://scholar.google.com/citations?user=ke8JkLgAAAAJ>