

Guanghui Qin

🌐 <https://gqin.me>

✉ guhqin@gmail.com

🎓 Google Scholar

Education

Johns Hopkins University

Ph.D. in Computer Science (Advisor: Benjamin Van Durme)

Aug 2019 – Aug 2024

Peking University

B.S. in Physics & Computer Science

Sept 2015 – Jun 2019

Experience

Microsoft Research Lab (MSR)

Washington, US

Research Intern (Mentor: Corby Rosset)

May 2023 – Aug 2023

Keywords: Large language model (LLM), Compressed text representation, Retrieval-augmented generation (RAG).

I researched efficient methods for long-context LLMs. I proposed a method to compress the context of LLaMA with a compression ratio of up to 20x with minimal performance tradeoffs. It worked on retrieval-augmented generation (RAG).

Microsoft Semantic Machines

Remote, US

Research Intern (Mentor: Anthony Platanios)

May 2022 – Aug 2022

Keywords: Dataset, Graph neural networks (GNNs), User action prediction.

I studied a new research problem for user action predictions. I built a 2TiB dataset from GitHub and implemented a GNN model to predict the user actions (e.g. commit and pull request).

Johns Hopkins University

Maryland, US

Visiting Researcher (Mentor: Hongyuan Mei and Jason Eisner)

Jun 2018 – Oct 2018

Keywords: Time-series models, Stochastic process, Datalog.

I worked on temporal event stream modeling. We proposed a particle smoothing solution to sample events from a neural Hawkes process. The stochastic process may interact with a deductive temporal database such as Datalog.

Microsoft Research-Asia (MSRA)

Beijing, China

Research Intern (Mentor: Jin-Ge Yao and Chin-Yew Lin)

Nov 2017 – Jun 2018

Keywords: Grounded language learning, Data-to-text generation.

I proposed a Semi-HMMs-based statistics model for grounding natural language to structured data, which can be used to induce templates for data-to-text generation.

Skills

- Programming languages: Python, Rust, JAVA, and C/C++. Other languages: Shell, \LaTeX , SQL.
- Experience in fine-tuning LLMs such as LLaMA, including distributed training, the use of LoRA.
- Experience data processing for biological/medical images: I developed a plugin for ImageJ to track objects using tools such as `scikit-image` and `scipy`.
- Machine learning tools: PyTorch, Lightning AI, DeepSpeed, FAISS, and PEFT.
- Network/Web: I have been hosting a proxy service (WallacePKU) since 2017. I implemented the proxy protocols and front-/back-end and maintain the database.

Awards

- Best Short Paper Awardee in NAACL
 - Outstanding Reviewer in EMNLP
 - Silver Medalist
- Association for Computational Linguistics (ACL), 2021
Association for Computational Linguistics (ACL), 2019
Chinese Physics Olympiad (CPhO), 2014

Selected Publications

- [Dodo: Dynamic Contextual Compression for Decoder-Only LMs.](#)
Guanghui Qin, Corby Rosset, Ethan C Chau, Nikhil Rao, and Benjamin Van Durme. In *Proceedings of Annual Meeting of the Association for Computational Linguistics (ACL, main)*. 2024.
- [Researchy Questions: A dataset of multi-perspective, compositional questions for LLM web agents.](#)
Corby Rosset, Ho-Lam Chung, **Guanghui Qin**, Ethan C Chau, Zhuo Feng, Ahmed Awadallah, Jennifer Neville, and Nikhil Rao. In *arXiv*. 2024.
- [Ras suppression potentiates actomyosin-driven polarization in migrating cells.](#)
Yiyan Lin*, Dhiman S Pal*, Parijat Banerjee, Tatsat Banerjee, **Guanghui Qin**, Yu Deng, Jane Borleis, Pablo A Iglesias, and Peter Devreotes. In *Nature Cell Biology (accepted, issue TBA)*. 2024.
- [Nugget: Neural Agglomerative Embeddings of Text.](#)
Guanghui Qin and Benjamin Van Durme. In *Proceedings of the Conference on International Conference on Machine Learning (ICML)*. 2023.
- [The NLP Task Effectiveness of Long-Range Transformers.](#)
Guanghui Qin, Yukun Feng, and Benjamin Van Durme. In *Proceedings of the European Chapter of the Association for Computational Linguistics (EACL, oral)*. 2023.
- [Learning How to Ask: Querying LMs with Mixtures of Soft Prompts.](#)
Guanghui Qin and Jason Eisner. In *Proceedings of Conference of the North American Chapter of the Association for Computational Linguistics (NAACL, short)*. 2021. **🏆 Best Short Paper Award**
- [Everything Is All It Takes: A Multipronged Strategy for Zero-Shot Cross-Lingual Information Extraction.](#)
Mahsa Yarmohammadi, Shijie Wu, Marc Marone, Haoran Xu, Seth Ebner, **Guanghui Qin**, Yunmo Chen, Jialiang Guo, Craig Harman, Kenton Murray, Aaron S White, Mark Dredze, and Benjamin Van Durme. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP, oral)*. 2021.
- [LOME: Large Ontology Multilingual Extraction.](#)
Patrick Xia*, **Guanghui Qin***, Siddharth Vashishtha, Yunmo Chen, Tongfei Chen, Chandler May, Craig Harman, Kyle Rawlins, Aaron S White, and Benjamin Van Durme. In *Proceedings of Conference of the European Chapter of the Association for Computational Linguistics (EACL, demo)*. 2021.
- [Iterative Paraphrastic Augmentation with Discriminative Span-based Alignment.](#)
Ryan Culkin, J Edward Hu, Elias Stengel-Eskin, **Guanghui Qin**, and Benjamin Van Durme. In *Transactions of the Association for Computational Linguistics (TACL)*, 9:494-509. 2021.
- [Neural Datalog Through Time: Informed Temporal Modeling via Logical Specification.](#)
Hongyuan Mei, **Guanghui Qin**, Minjie Xu, and Jason Eisner. In *Proceedings of the Conference on International Conference on Machine Learning (ICML, oral)*. 2020.
- [Imputing Missing Events in Continuous-Time Event Streams.](#)
Hongyuan Mei, **Guanghui Qin**, and Jason Eisner. In *Proceedings of the Conference on International Conference on Machine Learning (ICML, oral)*. 2019.
- [Learning Latent Semantic Annotations for Grounding Natural Language to Structured Data.](#)
Guanghui Qin, Jin-Ge Yao, Xuening Wang, Jinpeng Wang, and Chin-Yew Lin. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP, oral)*. 2018.
- [Data2Text Studio: Automated Text Generation from Structured Data.](#)
Longxu Dou, **Guanghui Qin**, Jinpeng Wang, Jin-Ge Yao, and Chin-Yew Lin. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP, demo)*. 2018.

Academic Service

I serve as a reviewer for the conferences of NeurIPS (2019 and 2020 as secondary; 2021 to 2023), ICLR (2019 and 2020 as secondary; 2021, 2023, and 2024), ICML (2020 and 2021), ACL (2021), EMNLP (2019 to 2022; *outstanding reviewer award* in 2019), NAACL (2024), AAAI (2021), and AKBC (2020 as secondary).

Last updated on May 16, 2024.